

STA414 Assignment 2

Q1 Implementing the model

(a) Implementing log_prior.

```
using Revise # lets you change A2funcs without restarting julia!
include("A2_src.jl")
using Plots
using Statistics: mean
using Zygote
using Test
using Logging
using .A2funcs: log1pexp # log(1 + exp(x)) stable
using .A2funcs: factorized_gaussian_log_density
using .A2funcs: skillcontour!
using .A2funcs: plot_line_equal_skill!

function log_prior(zs)
    return factorized_gaussian_log_density(0,0,zs)
end
```

(b) Implementing logp_a_beats_b.

```
function logp_a_beats_b(za,zb)
    return -log1pexp(zb-za)
end
```

(c) Implementing all_games_log_likelihood.

```
function all_games_log_likelihood(zs,games)
    zs_a = zs[games[:,1],:]
    zs_b = zs[games[:,2],:]
    likelihood = logp_a_beats_b.(zs_a, zs_b)
    return sum(likelihood, dims = 1)
end
```

(d) Implementing joint_log_density.

```
function joint_log_density(zs,games)
    return sum(log_prior(zs) .+ all_games_log_likelihood(zs, games), dims=1)
end
```

```
35 @testset "Test shapes of batches for likelihoods" begin
36     B = 15 # number of elements in batch
37     N = 4 # Total Number of Players
38     test_zs = randn(4,15)
39     test_games = [1 2; 3 1; 4 2] # 1 beat 2, 3 beat 1, 4 beat 2
40     @test size(test_zs) == (N,B)
41     #batch of priors
42     @test size(log_prior(test_zs)) == (1,B)
43     # loglikelihood of p1 beat p2 for first sample in batch
44     @test size(logp_a_beats_b(test_zs[1,1],test_zs[2,1])) == ()
45     # loglikelihood of p1 beat p2 broadcasted over whole batch
46     @test size(logp_a_beats_b.(test_zs[1,:],test_zs[2,:])) == (B,)
47     # batch loglikelihood for evidence
48     @test size(all_games_log_likelihood(test_zs,test_games)) == (1,B)
49     # batch loglikelihood under joint of evidence and prior
50     @test size(joint_log_density(test_zs,test_games)) == (1,B)
51 end |> Test.DefaultTestSet
52
```

REPL


Test Summary:	Pass	Total
Test shapes of batches for likelihoods	6	6

Q2 Examining the posterior for only two players and toy data

(a) plot the log prior graph.

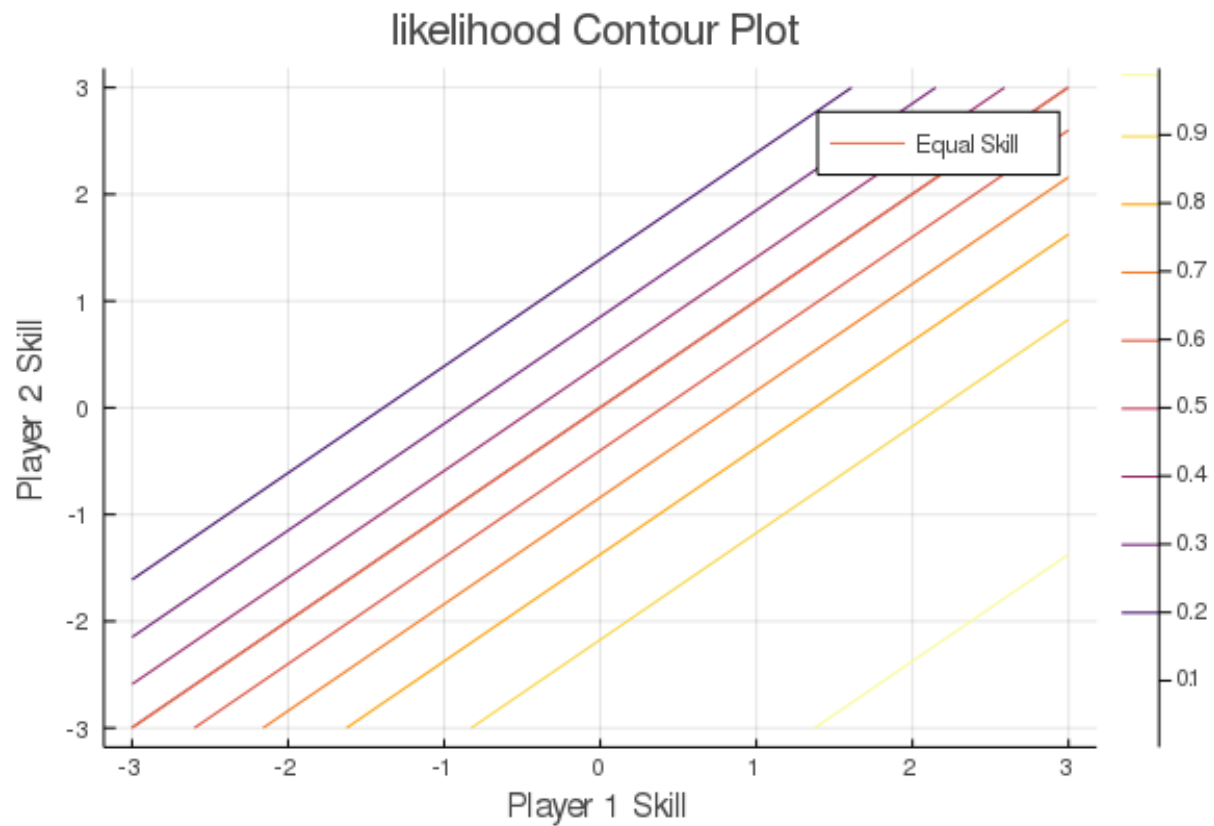
```
# Q2(a)
plot(title="logprior Contour Plot",
      xlabel = "Player 1 Skill",
      ylabel = "Player 2 Skill"
    )

log_p(zs) = exp(log_prior(zs))
skillcontour!(log_p)
plot_line_equal_skill!()
savefig(joinpath("plots","prior.png"))
```




(b) Plot the log likelihood Contour Plot.

```
plot(title="likelihood Contour Plot",  
      xlabel = "Player 1 Skill",  
      ylabel = "Player 2 Skill"  
    )  
skillcontour!(zs -> exp.(logp_a_beats_b(zs[1], zs[2])))  
plot_line_equal_skill!()  
savefig(joinpath("plots","likelihood.png"))
```




(c) Plot the Contour of A B 1:0.

```
# Q2(c)
plot(title="A B 1:0 Contour Plot",
      xlabel = "Player 1 Skill",
      ylabel = "Player 2 Skill"
    )
skillcontour!(zs -> exp.(joint_log_density(zs, two_player_toy_games(1,0))))
plot_line_equal_skill!()
savefig(joinpath("plots","contour_a1b0.png"))
```




(d) Plot Contour of A B 10:0.

```
plot(title="A B 10:0 Contour Plot",
      xlabel = "Player 1 Skill",
      ylabel = "Player 2 Skill"
    )
skillcontour!(zs -> exp.(joint_log_density(zs, two_player_toy_games(10,0))))
plot_line_equal_skill!()
savefig(joinpath("plots","toy_svi_a10b0.png"))
```



(e) Plot Contour of A B 10:10.

```
plot(title="A B 10:10 Contour Plot",  
      xlabel = "Player 1 Skill",  
      ylabel = "Player 2 Skill"  
    )  
skillcontour!(zs -> exp.(joint_log_density(zs, two_player_toy_games(10,10))))  
plot_line_equal_skill!()  
savefig(joinpath("plots", "toy_svi_a10b10.png"))
```



Q3 Stochastic Variational Inference on Two Players and Toy Data

(a) Implement elbo.

```
function elbo(params, logp, num_samples)
    samples = exp.(params[2]) .* randn(size(params[1])[1], num_samples) .+ params[1]
    logp_estimate = logp(samples)
    logq_estimate = factorized_gaussian_log_density(params[1], params[2], samples)
    return sum(logp_estimate - logq_estimate) / num_samples
end
```

(b) Implement neg_toy_elbo.

```
function neg_toy_elbo(params; games = two_player_toy_games(1,0), num_samples = 100)
    # TODO: Write a function that takes parameters for q,
    # evidence as an array of game outcomes,
    # and returns the -elbo estimate with num_samples many samples from q
    logp(zs) = joint_log_density(zs, games)
    return -elbo(params, logp, num_samples)
end
```

(c) Implement fit_toy_variational_dist.

```
function fit_toy_variational_dist(init_params, toy_evidence; num_itrs=200,
    lr= 1e-2, num_q_samples = 10)
    params_cur = init_params
    for i in 1:num_itrs
        grad_params = gradient(params -> neg_toy_elbo(params; games = toy_evidence,
            num_samples = num_q_samples), params_cur)[1]
        mu_new = params_cur[1] - lr * grad_params[1]
        ls_new = params_cur[2] - lr * grad_params[2]
        params_cur = (mu_new, ls_new)
        # @info "neg_elbo: $(neg_toy_elbo(params_cur; games = toy_evidence,
        # num_samples = num_q_samples))"
        plot(title="fit toy variational distribution", xlabel = "Player 1 Skill",
            ylabel = "Player 2 Skill");

        target_post(zs) = exp.(joint_log_density(zs, toy_evidence))
        skillcontour!(target_post, colour=:red)
        plot_line_equal_skill!()
        mu = params_cur[1]
        logsig = params_cur[2]
        var_log_prior(zs) = factorized_gaussian_log_density(mu, logsig, zs)
        var_post(zs) = exp.(var_log_prior(zs))
        display(skillcontour!(var_post, colour=:blue))
    end
    return params_cur
end
```

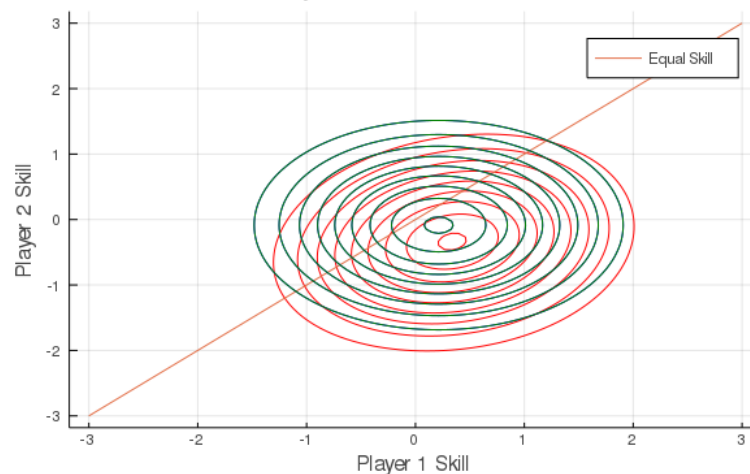
(d) Optimized variational approximation contours Plot of A B 1:0.

```
#Q1(d) 1:0
plot(title="A B 1:0 Toy SVI",
      xlabel = "Player 1 Skill",
      ylabel = "Player 2 Skill"
    )
# Plot the original graph
target_post(zs) = exp.(joint_log_density(zs, two_player_toy_games(1,0)))
skillcontour!(target_post, colour=:red)
plot_line_equal_skill!()

# Plot SVI
num_q_samples = 10
data = two_player_toy_games(1,0)
opt_params = fit_toy_variational_dist(toy_params_init, two_player_toy_games(1,0), num_itrs=200, lr= 1e-2, num_q_samples = 10)
println("neg_elbo (final loss): $(neg_toy_elbo(opt_params; games = two_player_toy_games(1,0), num_samples = num_q_samples))")
mu = opt_params[1]
logsig = opt_params[2]
var_log_prior(zs) = factorized_gaussian_log_density(mu, logsig, zs)
var_post(zs) = exp.(var_log_prior(zs))
display(skillcontour!(var_post, colour=:green))
savefig(joinpath("plots", "toy_svi_a1b0.png"))
```

neg_elbo (final loss): 0.7439268493244543

fit toy variational distribution



(e) Optimized variational approximation contours Plot of A B 10:0.

```
plot(title="A B 10:0 Toy SVI",
      xlabel = "Player 1 Skill",
      ylabel = "Player 2 Skill"
    )
target_post(zs) = exp.(joint_log_density(zs, two_player_toy_games(10,0)))
skillcontour!(target_post, colour=:red)
plot_line_equal_skill!()

# Plot the new graph
opt_params = fit_toy_variational_dist(toy_params_init, two_player_toy_games(10,0), num_itrs=200, lr= 1e-2, num_q_samples = 10)
num_q_samples = 10
println("neg_elbo (final loss): $(neg_toy_elbo(opt_params; games = two_player_toy_games(10,0), num_samples = num_q_samples))")
mu = opt_params[1]
logsig = opt_params[2]
var_log_prior(zs) = factorized_gaussian_log_density(mu, logsig, zs)
var_post(zs) = exp.(var_log_prior(zs))
display(skillcontour!(var_post, colour=:green))
savefig(joinpath("plots", "toy_svi_a10b0.png"))
```

neg_elbo (final loss): 2.9131816235518913

(f) Optimized variational approximation contours Plot of A B 10:10


```
opt_params = fit_toy_variational_dist(toy_params_init, two_player_toy_games(10,10); num_itrs=200, lr= 1e-2, num_q_samples = 10)
num_q_samples = 10
println("neg_elbo (final loss): $(neg_toy_elbo(opt_params; games = two_player_toy_games(10,10), num_samples = num_q_samples))")

plot(title="Toy SVI with A:B = 10:10",
      xlabel = "Player 1 Skill",
      ylabel = "Player 2 Skill"
    )
target_post(zs) = exp.(joint_log_density(zs, two_player_toy_games(10,10)))
skillcontour!(target_post, colour=:red)
plot_line_equal_skill!()

mu = opt_params[1]
logsig = opt_params[2]
var_log_prior(zs) = factorized_gaussian_log_density(mu, logsig, zs)
var_post(zs) = exp.(var_log_prior(zs))
display(skillcontour!(var_post, colour=:green))
savefig(joinpath("plots", "toy_svi_a10b10.png"))
```

neg elbo (final loss): 15.746682651790952

Toy SVI with A:B = 10:10



Q4 Approximate inference conditioned on real data

(a) Yes.

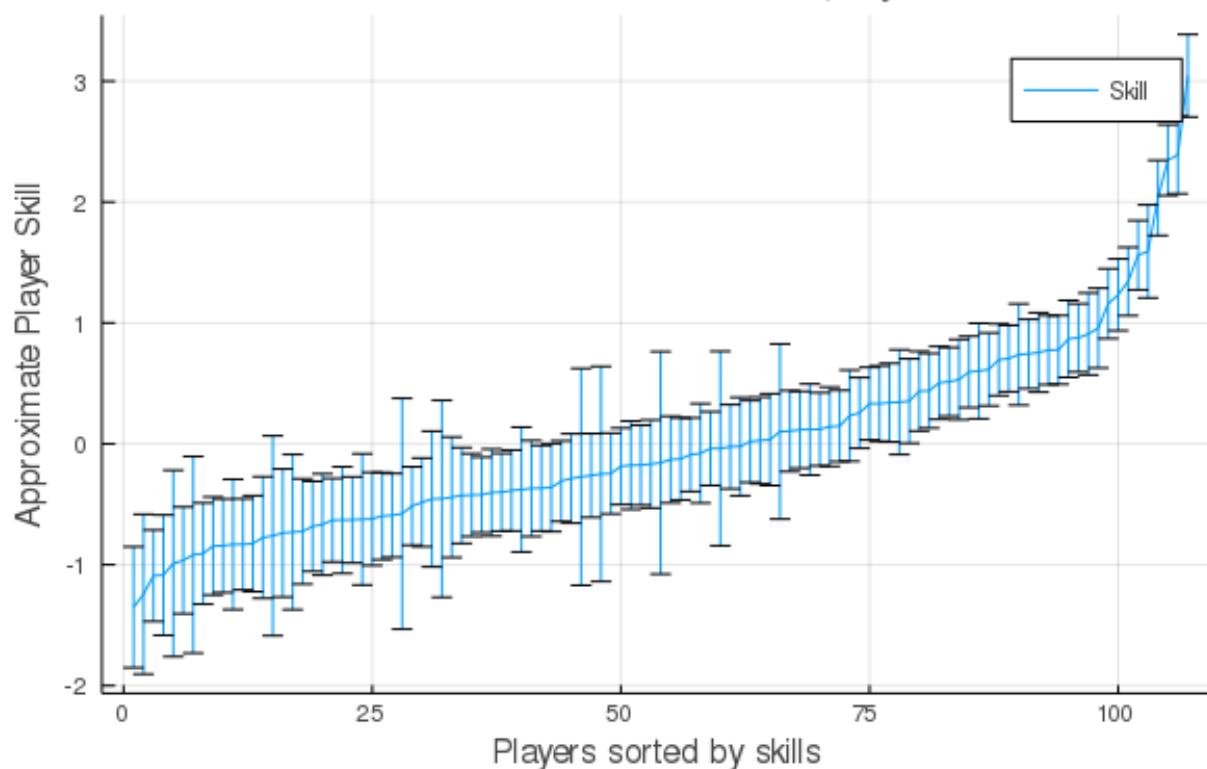
(b) loss: 1185.287000374451

```
function fit_variational_dist(init_params, tennis_games; num_itrs=200, lr= 1e-2, num_q_samples = 10)
    params_cur = init_params
    for i in 1:num_itrs
        grad_params = gradient(params -> neg_toy_elbo(params; games = tennis_games, num_samples = num_q_samples), params_cur)[1]
        mu_new = params_cur[1] - lr * grad_params[1]
        ls_new = params_cur[2] - lr * grad_params[2]
        params_cur = (mu_new, ls_new)
        @info "neg_elbo: $(neg_toy_elbo(params_cur; games = tennis_games, num_samples = num_q_samples))"
    end
    println("neg_elbo (final loss): $(neg_toy_elbo(params_cur; games = tennis_games, num_samples = num_q_samples))")
    return params_cur
end | > fit_variational_dist
```

(c)

```
perm = sortperm(means)
plot(title="Mean and variance of all players",
     xlabel = "Players sorted by skills",
     ylabel = "Approximate Player Skill"
)
plot!(means[perm], yerror=exp.(logstd[perm]), label="Skill")
savefig(joinpath("plots", "player_mean_var.png"))
```

Mean and variance of all players



(d)

```
desc_perm = sortperm(means, rev=true)
print("Top 10 players are: ")
for i in 1:10
    print(player_names[desc_perm][i], "\n")
end
```

Top 10 players are: Novak-Djokovic
Roger-Federer
Rafael-Nadal
Andy-Murray
Robin-Soderling
David-Ferrer
Jo-Wilfried-Tsonga
Tomas-Berdych
Juan-Martin-Del-Potro
Richard-Gasquet

(f)

$$\begin{aligned} \text{f)} \quad A \begin{pmatrix} z_A \\ z_B \end{pmatrix} &= \begin{pmatrix} y_A \\ y_B \end{pmatrix} \\ \text{solve for } A, A &= \begin{pmatrix} 1 & -1 \\ 0 & 1 \end{pmatrix} \\ \text{Since } \begin{pmatrix} z_A \\ z_B \end{pmatrix} &\sim N \left(\begin{pmatrix} \mu_A \\ \mu_B \end{pmatrix}, \begin{pmatrix} \sigma_A^2 & 0 \\ 0 & \sigma_B^2 \end{pmatrix} \right) = N(\mu, \Sigma) \\ \Rightarrow \begin{pmatrix} y_A \\ y_B \end{pmatrix} &\sim N \left(A\mu, A\Sigma A' \right) = N \left(\begin{pmatrix} \mu_A/\sigma_B \\ \mu_B \end{pmatrix}, \begin{pmatrix} \sigma_A^2 + \sigma_B^2 & -\sigma_B^2 \\ -\sigma_B^2 & \sigma_B^2 \end{pmatrix} \right) \\ \Rightarrow P(y_A | x) &\sim N(\mu_A/\sigma_B, \sigma_A^2 + \sigma_B^2) \\ \Rightarrow \text{let } \mu_{ja} &= N(\mu_A/\sigma_B, \sigma_A^2 + \sigma_B^2) \\ P(y_A > 0) &= 1 - \text{cdf}(\mu_{ja}, 0) \end{aligned}$$

(g)

```
# Q4(g)
# Exact approach
exact_rf_better = 1 - cdf(Normal(0,1), (mu_RN - mu_RF)/sqrt(var_RF + var_RN))
# SM approach
MC_size = 10000
samples_RF = randn(MC_size) * exp(-1.1412578223971224) .+ mu_RF
samples_RN = randn(MC_size) * exp(-1.2163985016732244) .+ mu_RN
MC_rf_better = count(x->x==1, samples_RF .> samples_RN) / MC_size
...
The result of SM of RF is better than RN is: 0.5755
```

(h)

```
# Q4(h)
lowest_idx = desc_perm[num_players]
mu_lowest = means[lowest_idx]
var_lowest = exp(logstd[lowest_idx])^2
# Exact prob
exact_rf_better_than_worst = 1 - cdf(Normal(0,1), (mu_lowest - mu_RF)/sqrt(var_RF + var_lowest))
# SM approach
samples_lowest = randn(MC_size) * exp(logstd[lowest_idx]) .+ mu_lowest
MC_h = count(x->x==1,samples_RF .> samples_lowest) / MC_size
```

The result of SM of RF is better the worst is: 1.0

(i)

(b), (c), and (e)

Jinhan Mei
1003137937

Jinhan Mei
1003137937

